

Generative Modeling using Matchgate Circuits

Peter Luo, Hong Ye Hu, Susanne F. Yelin


September 25, 2023

Neural Networks

- Neural Networks are a means by which machine learning is carried out, and its architecture is inspired by neurons

A simple neural network


input layer hidden layer output layer



- Each neuron will output


$$f\left(b + \sum w_i x_i\right),$$

where f is an “activation function”, the w_i are the “weights”, the b is the “bias”, and the x_i being the inputs



How Neural Networks Learn


- Define a “**loss function**” $L(\mathbf{w}, \mathbf{b})$ which captures how close the network’s output is to the desired output



- We can use optimization algorithms to adjust \mathbf{w} and \mathbf{b} to get closer and closer to the minimum of L


Generative Modeling

- Rather than doing a classification or prediction task, generative models strive to generate more instances of their training data
- Examples: DALL-E, Voice Generators
- **Goal of this project:** To use a Matchgate circuit and the MNIST dataset to generate new images of handwritten digits



Generative Adversarial Networks (GAN)

- Consists of two networks G and D with parameters θ and ϕ and their own loss functions $L_G(\theta, \phi)$ and $L_D(\theta, \phi)$
- $G : \{0, 1\}^{28 \times 28} \rightarrow \{0, 1\}^{28 \times 28}$ and $D : \{0, 1\}^{28 \times 28} \rightarrow [0, 1]$



- Assume that the real data $\mathbf{x}_1, \dots, \mathbf{x}_n$ comes from an underlying probability distribution p_{real} , and recall that the input to G are samples $\mathbf{z}_1, \dots, \mathbf{z}_n$ from p_{prior}
- The loss functions are

$$L_G(\boldsymbol{\theta}, \phi) = - \mathbb{E}_{\mathbf{z} \sim p_{\text{prior}}} [\log(D(G(\mathbf{z})))]$$

$$L_D(\boldsymbol{\theta}, \phi) = - \left(\mathbb{E}_{\mathbf{x} \sim p_{\text{real}}} [\log(D(\mathbf{x}))] + \mathbb{E}_{\mathbf{z} \sim p_{\text{prior}}} [\log(1 - D(G(\mathbf{z})))] \right).$$



Quantum Computing

- Leverages principles from quantum mechanics to do computations
- Consequently, it can solve certain problems exponentially faster than classical computers, reducing runtime from millions of years to just a few minutes
- It can also break some commonly used encryption methods, such as RSA or ECC
- It can simulate molecular interactions efficiently, which has implications in drug discovery and the material sciences

Limitations

- Quantum computers need to be kept at very cold temperatures, thus requiring a large amount of energy
- Real quantum devices are noisy
- Number of qubits that we're able to run on the best quantum computers is relatively small (on the order of 10^2)
- A lot is known about what is **theoretically possible**, but not a lot is known about their actual performance
- Temporary solution: classical simulation

Quantum Computing

Classical computing:

- Consists of logic gates (e.g. AND, NOT, OR)
- One classical bit: a 0 or a 1

Quantum computing:

- Consists of quantum gates, represented by unitary matrices with complex entries
- One quantum bit: a vector of two complex numbers, which describes the probability of being measured in the 0 or 1 state

Quantum Bits

- One qubit looks like

$$a|0\rangle + b|1\rangle \rightarrow \begin{pmatrix} a \\ b \end{pmatrix}$$

- $|a|^2$ and $|b|^2$ are the probabilities so we require $|a|^2 + |b|^2 = 1$
- For two qubits,

$$a|00\rangle + b|01\rangle + c|10\rangle + d|11\rangle \rightarrow \begin{pmatrix} a \\ b \\ c \\ d \end{pmatrix}$$

- The number of complex numbers we need to keep track of scales **exponentially** with the number of qubits
- Quantum computers are hard to simulate classically

Quantum Gates

- An n -qubit gate is a $2^n \times 2^n$ unitary matrix with complex entries
- A matrix U is unitary if $UU^\dagger = I$, where \dagger means the conjugate transpose
- 1-qubit examples:

$$H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}, \quad \begin{pmatrix} 1 & 0 \\ 0 & i \end{pmatrix}, \quad \begin{pmatrix} \cos(\theta/2) & -\sin(\theta/2) \\ \sin(\theta/2) & \cos(\theta/2) \end{pmatrix}$$


- Example:

$$H|0\rangle = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix} \begin{pmatrix} 1 \\ 0 \end{pmatrix} = \begin{pmatrix} 1/\sqrt{2} \\ 1/\sqrt{2} \end{pmatrix} \rightarrow \frac{1}{\sqrt{2}}|0\rangle + \frac{1}{\sqrt{2}}|1\rangle$$

- To simulate quantum computers is to do matrix multiplications

Parametrized Quantum Circuit

- Quantum gates can be **parametrized**
- By replacing the neural network with a parametrized quantum circuit and then measuring its output, we obtain a **hybrid** quantum-classical system compatible with machine learning procedures



Computing Gradients

- Parameter-Shift Rule
- If $L(\theta)$ is the expectation value of some quantity with respect to the output of a quantum circuit with gates of the form $U_G(\theta) = e^{-ia\theta G}$ where $G^2 = I$, then

$$\frac{\partial L}{\partial \theta_i} = r \left(L \left(\theta + \frac{\pi}{4r} \vec{e}_i \right) + L \left(\theta - \frac{\pi}{4r} \vec{e}_i \right) \right)$$

- Proof uses

$$U_G(\theta) = e^{-i\theta G} = I \cos \theta - iG \sin \theta$$

and

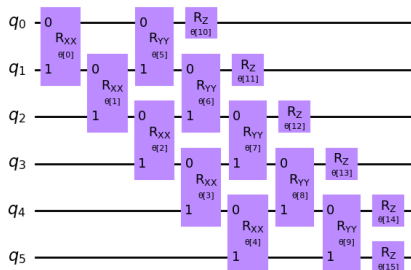
$$\frac{d}{d\theta} U_G(\theta) = -iG e^{-i\theta G} = -iG U_G(\theta)$$

Matchgates

- Matchgates are a restricted class of parametrized quantum gates, generated by

$$e^{-iX \otimes X \theta/2}, \quad e^{-iY \otimes Y \theta/2}, \quad e^{-iX \otimes Y \theta/2}, \quad e^{-iY \otimes X \theta/2}, \quad e^{-iZ \theta/2}$$

- Matchgates are **differentiable** with respect to θ and can be simulated in **polynomial time**



Valiant's Definition

- Matchgates were first introduced by Valiant (2001)
- For a given undirected weighted graph G where $V(G) = [n]$, define the **Pfaffian** to be

$$\text{Pf}(B) = \sum_{\substack{\text{perfect matchings} \\ M}} \text{sgn}(M) \prod_{(ij) \in M} w_{ij},$$

where B is the adjacency matrix of G

- Theorem (Cayley): $\text{Pf}(B)^2 = \text{Det}(B)$
- Pfaffians are computable in polynomial time

- Define the **Pfaffian sum** to be

$$\text{PfS}(B) = \sum_{A \subseteq [n]} \left(\prod_{i \in A} \lambda_i \right) \text{Pf}(B[A])$$

- Pfaffian sums are just Pfaffians:

$$\text{PfS}(B) = \begin{cases} \text{Pf}(B + \Lambda_n) & n \text{ even} \\ \text{Pf}(B^+ + \Lambda_{n+1}) & n \text{ odd} \end{cases}$$

where

$$\Lambda_n = \begin{pmatrix} 0 & \lambda_1 \lambda_2 & -\lambda_1 \lambda_3 & \cdots \\ -\lambda_1 \lambda_2 & 0 & \lambda_2 \lambda_3 & \cdots \\ \lambda_1 \lambda_3 & -\lambda_2 \lambda_3 & \ddots & \\ \vdots & \vdots & & \end{pmatrix}$$

- A **matchgate** Γ is (G, X, Y, T) where G is a graph and X, Y, T partition the vertices. For $Z \subseteq X \cup Y$, define

$$\chi(\Gamma, Z) = \mu(\Gamma, Z) \text{PfS}(G - Z)$$

- The **character matrix** $\chi(\Gamma)$ is the $2^{|X|} \times 2^{|Y|}$ matrix with

$$\chi(\Gamma)_{i,j} = \chi(\Gamma, X_i \cup Y_j).$$

- Defined in this way, not all matchgates are unitary
- Take the ones that are unitary and interpret them as quantum gates
- For a Matchgate circuit C ,

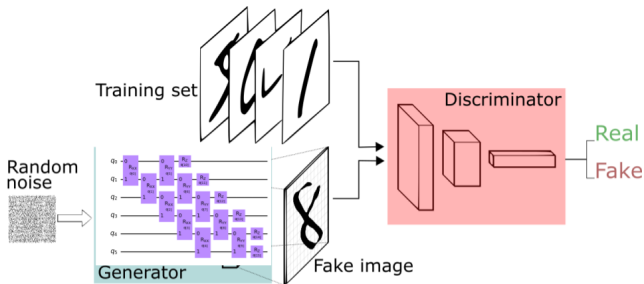
$$|\langle \underline{y} | C | \underline{x} \rangle|^2 = |\text{PfS}(M_{\underline{x}, \underline{y}})|^2.$$

Matchgate Speedup

- Valiant: Computing Pfaffians is fast
- Terhal and DiVincenzo (2001): There is a correspondence to noninteracting fermions in 1D, and computing determinants is fast

Setup

- 1 Use a parametrized quantum circuit made of matchgates as our G
- 2 Use a classical neural network as D
- 3 Train it just like you would train a GAN, i.e. use the same loss functions



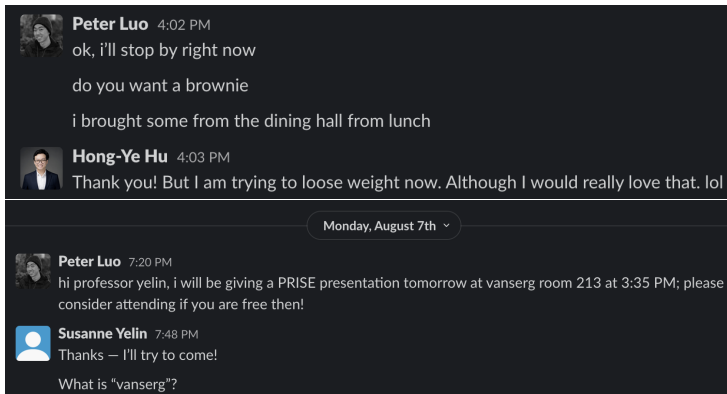
- Goal: Understand the performance of QML on a real dataset with a large number of qubits

Code

- Written in Julia, using Yao package
- Utilized FLOYao package to harness Matchgate optimizations
- Current progress

Acknowledgements

- Hong-Ye Hu and Susanne Yelin



The screenshot shows a WhatsApp chat interface with a dark background. At the top, there is a header bar with the name "Peter Luo" and the time "4:02 PM". Below this, a message from Peter Luo reads: "ok, i'll stop by right now", "do you want a brownie", and "i brought some from the dining hall from lunch". The next message is from Hong-Ye Hu, with a header bar showing "Hong-Ye Hu" and "4:03 PM". The message from Hong-Ye Hu says: "Thank you! But I am trying to loose weight now. Although I would really love that. lol". Below this message is a separator bar with the date "Monday, August 7th" and a dropdown arrow. The next message is from Peter Luo, with a header bar showing "Peter Luo" and "7:20 PM". The message from Peter Luo reads: "hi professor yelin, i will be giving a PRISE presentation tomorrow at vanserg room 213 at 3:35 PM; please consider attending if you are free then!". The final message is from Susanne Yelin, with a header bar showing "Susanne Yelin" and "7:48 PM". The message from Susanne Yelin reads: "Thanks — I'll try to come!" and "What is 'vanserg'?".

Peter Luo 4:02 PM
ok, i'll stop by right now
do you want a brownie
i brought some from the dining hall from lunch

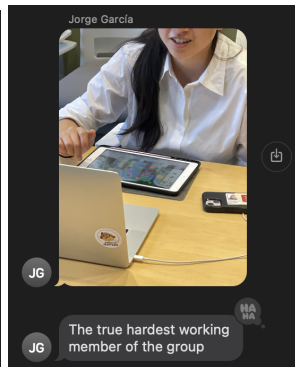


Hong-Ye Hu 4:03 PM
Thank you! But I am trying to loose weight now. Although I would really love that. lol

Monday, August 7th ▾

Peter Luo 7:20 PM
hi professor yelin, i will be giving a PRISE presentation tomorrow at vanserg room 213 at 3:35 PM; please consider attending if you are free then!

Susanne Yelin 7:48 PM
Thanks — I'll try to come!
What is "vanserg"?

- Yelin group undergrads: Jorge, Fiona, Victoria, and Kevin



Questions?

